

Kubernetes Networking for 'Stackers

Karthik Prabhakar Tigera kp@tigera.io

### Today

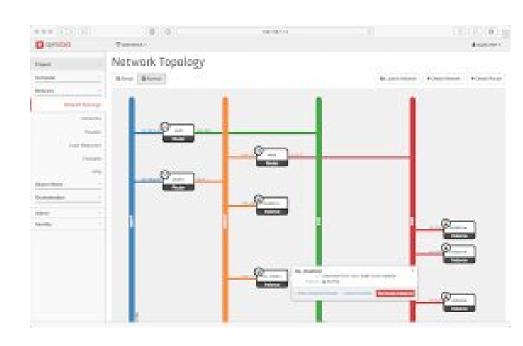
- 1. How did we get here? (a.k.a. Neutron's Design Thinking)
- 2. Where are we headed? (a.k.a. Kubernetes Network Design Thinking)
- 3. Kubernetes Network concepts
- 4. Live Demo: Kubernetes networking can be simple

### Tomorrow (Wed): Hynes Level 3 - Ballroom B - 4:30pm

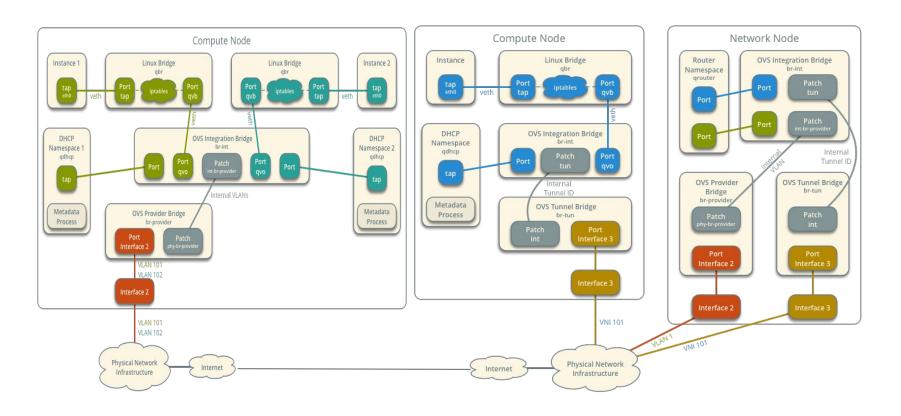
- A. <u>Simple</u>, <u>Scalable</u> Approaches to deploy Kubernetes with OpenStack
  - a. Kubernetes alongside OpenStack
  - b. Kubernetes in OpenStack
  - c. OpenStack in Kubernetes
- B. AT&T Overview & Demo: Deploying OpenStack on Kubernetes (with OpenStack-Helm and Calico)

### Remnants from the early days of VM Networking

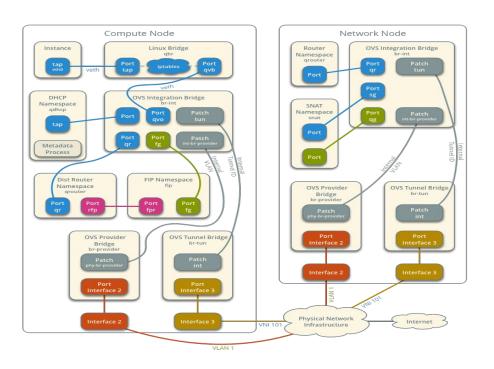
- => Conflate isolation with network topology
- => Complex SDN controllers to layer overlay spaghetti onto physical infra
- => Random mesh of bridges, vswitches, tunnels, L3 backhauls, and security enforcement points
- => Coarse security group rules decoupled from dynamic workloads

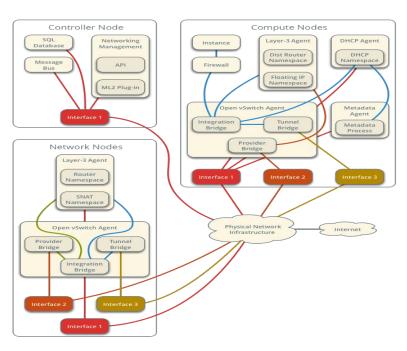


### **Exhibit A:** Neutron with OpenvSwitch (L2)

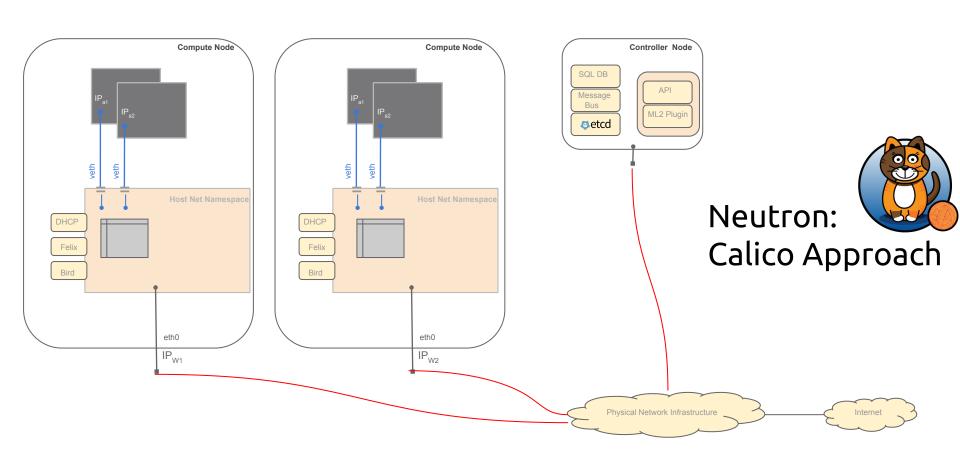


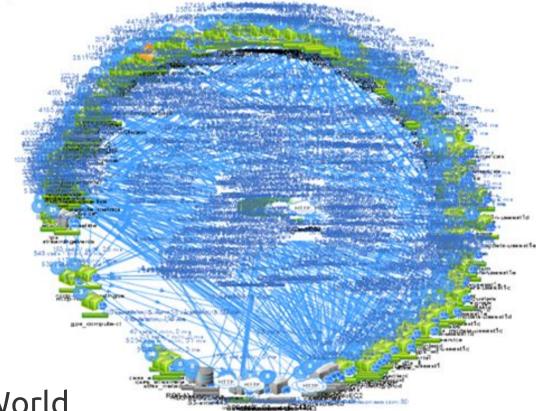
### Exhibit B: Neutron with OpenvSwitch and DVR/VRRP











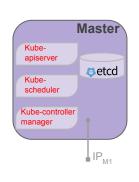
The Dynamic World of Microservices

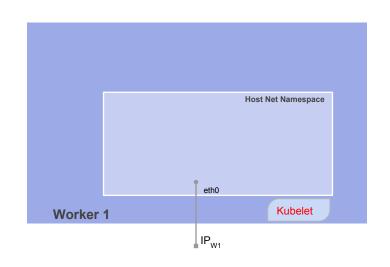


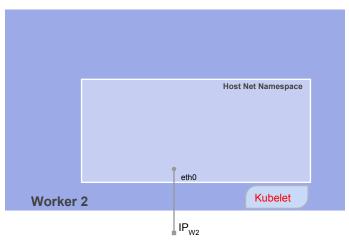
### The Kubernetes Approach

- => Assume **IP connectivity** between nodes within cluster
- => IP per Pod
- => CNI plugin abstraction enabling choice of network plugin for endpoint traffic
- => <u>Declarative</u> isolation using rich **Network Policy** constructs
- => "Services" map to pods dynamically
- => Optional **Ingress** controllers and resources
- => Choice of **Service Discovery** options, including KubeDns

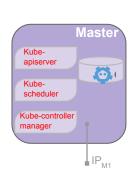
### The Kubernetes Approach: Nodes

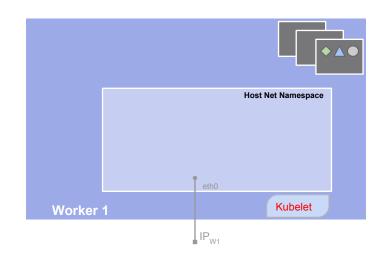


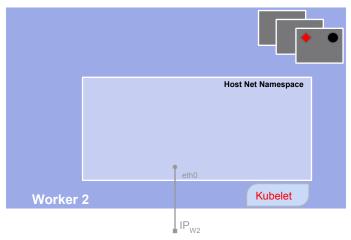




### The Kubernetes Approach: Pods



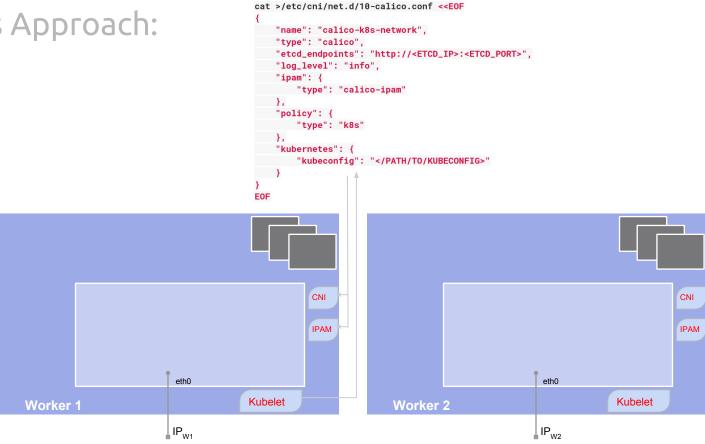


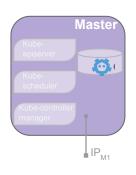


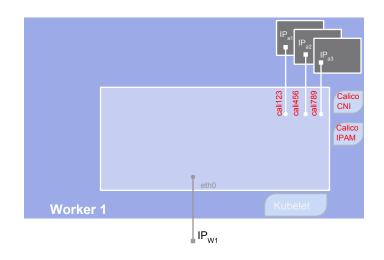
### The Kubernetes Approach: CNI and IPAM

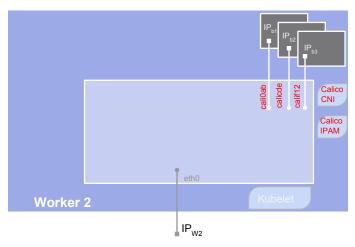


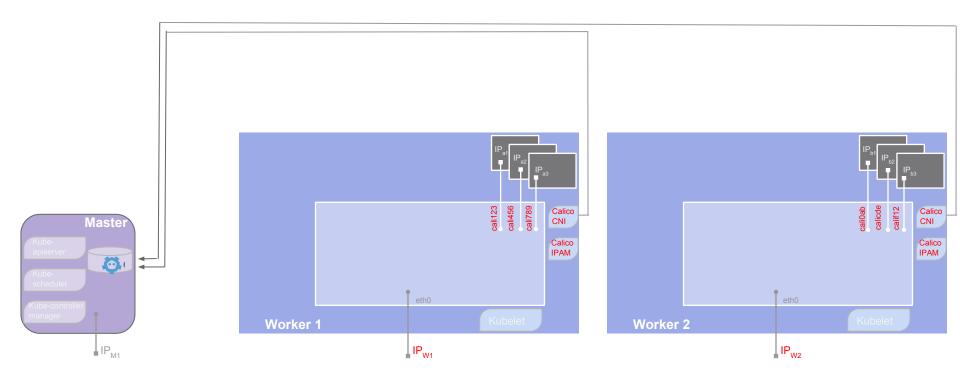
IP<sub>M1</sub>

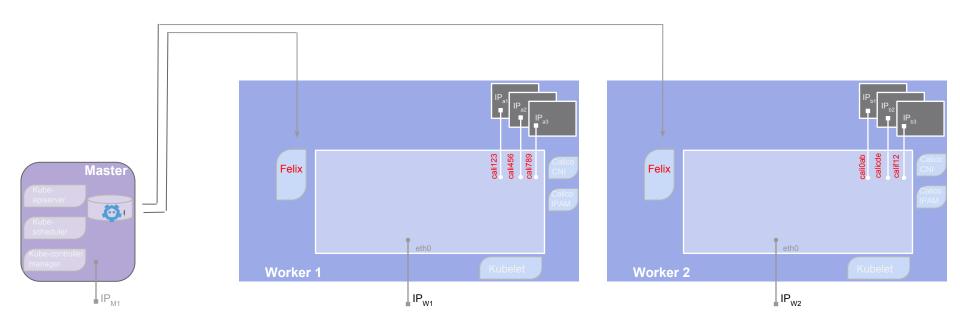


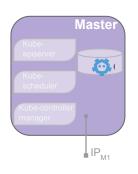


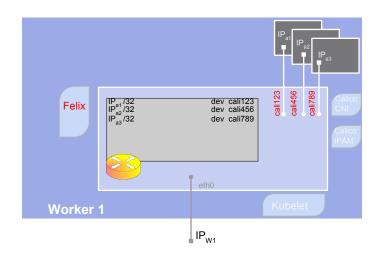


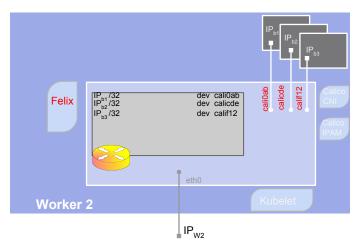


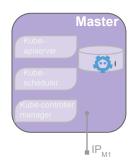


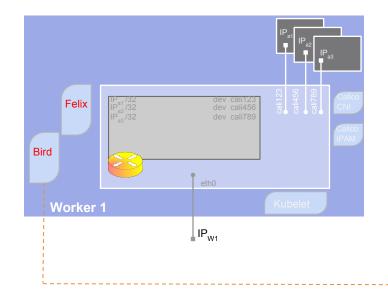


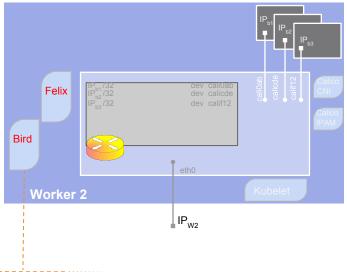


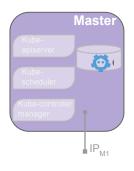


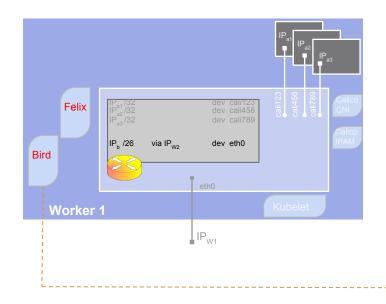


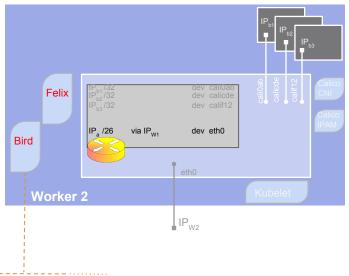




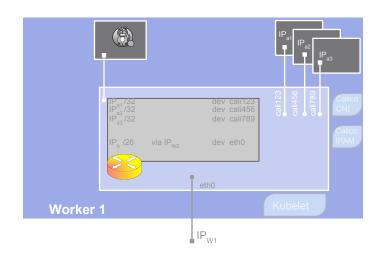


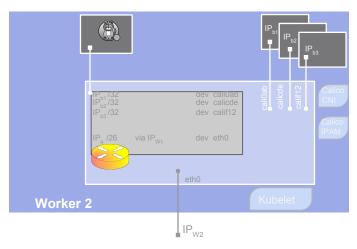


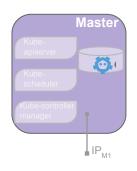


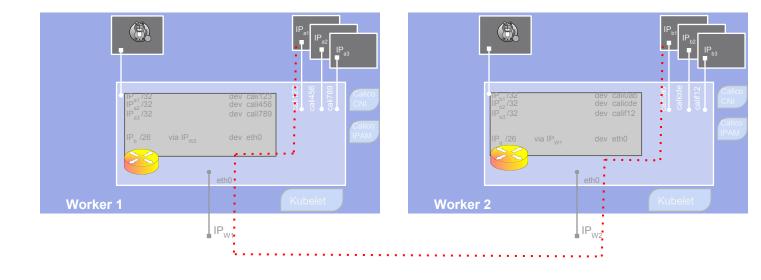






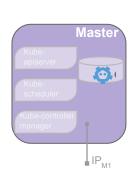


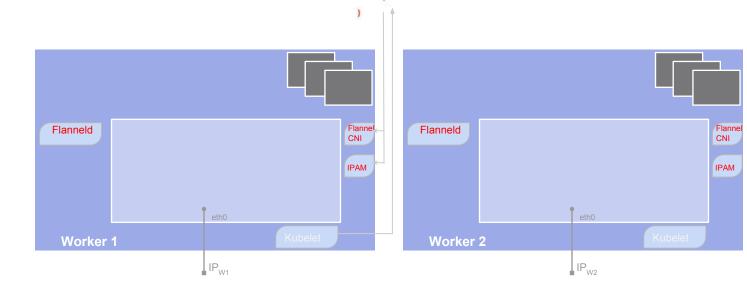




The Kubernetes Approach: # flannel (Pod-Pod Networking)





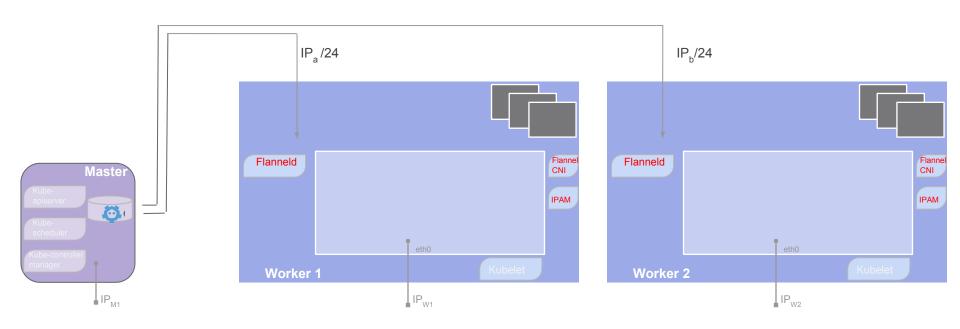


"name": "mynet",

"type": "bridge", "mtu": 1472, "ipMasq": false, "isGateway": true, "ipam": {

> "type": "host-local", "subnet": "10.1.17.0/24"

# The Kubernetes Approach: # flannel (Pod-Pod Networking)



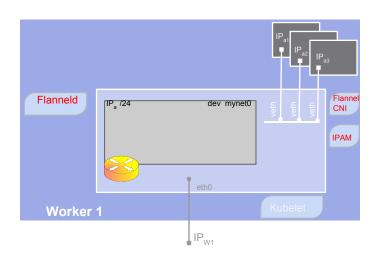
## The Kubernetes Approach: ff flannel (Pod-Pod Networking)

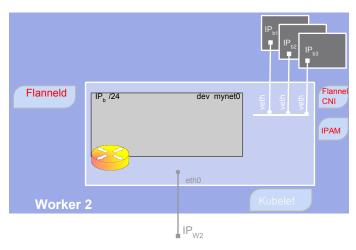
#### Flannel Backends:

UDP
VXLAN
Host-gateway
AWS-VPC
GCE
Ali-VPC

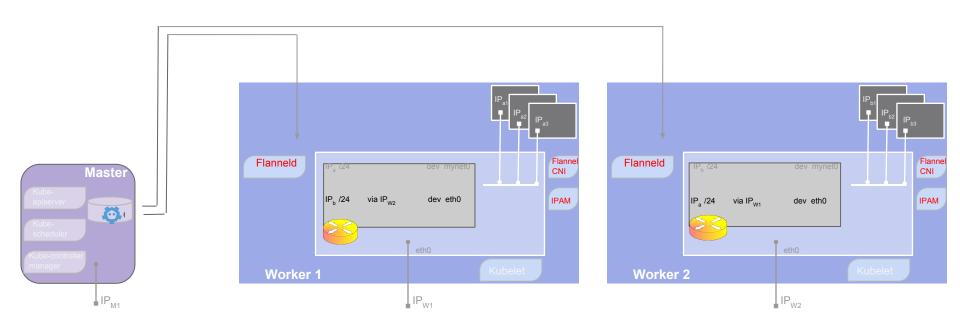
Alloc



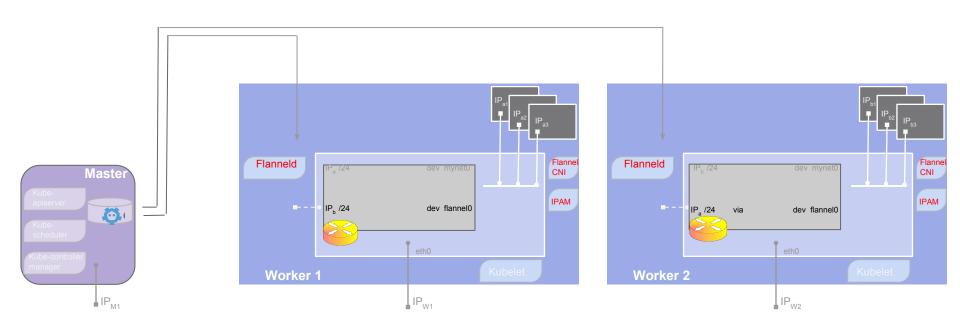




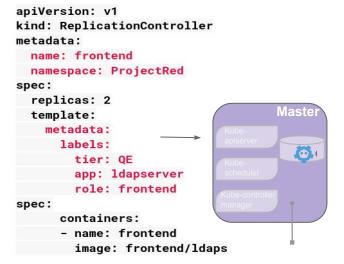
The Kubernetes Approach: # flannel: Host Gateway mode (Pod-Pod Networking)

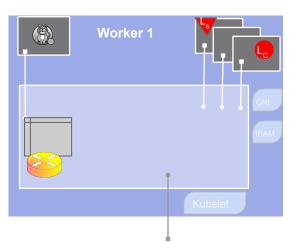


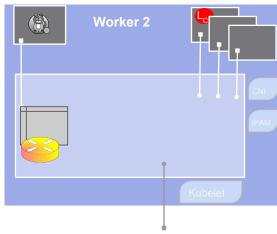
The Kubernetes Approach: ff flannel: VXLAN mode (Pod-Pod Networking)



## The Kubernetes Approach: Network Policy Namespaces, Labels and Selectors



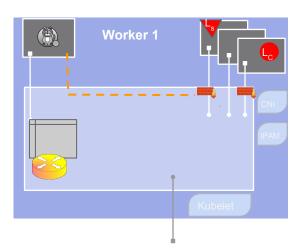


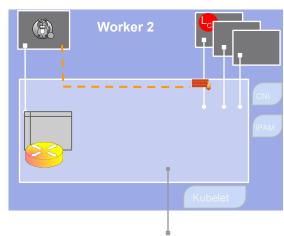




### The Kubernetes Approach: Network Policy

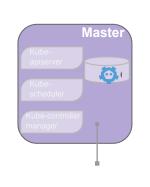
```
kind: NetworkPolicy
apiVersion: extensions/v1beta1
metadata:
 namespace: ProjectRed
 name: ldap-frontend
spec:
 podSelector:
    matchLabels:
      tier: OE
      role: frontend
      app: ldapserver
                                         Master
    - from:
  ingress:
                                            O.
    - from:
        - podSelector:
            matchLabels:
              tier: OE
              role: frontend
              app: ldapclient
      ports:
        - protocol: TCP
          port: 636
```

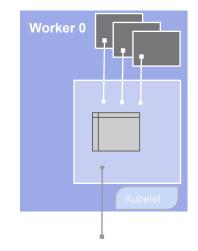


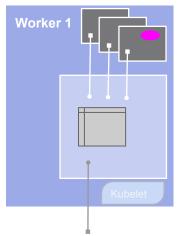


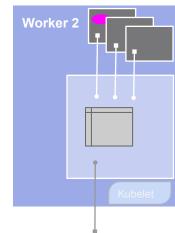


### The Kubernetes Approach: Services

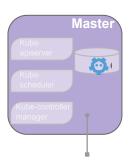


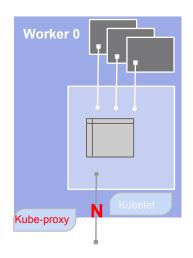


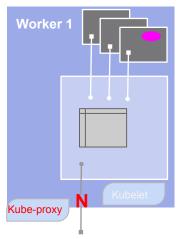


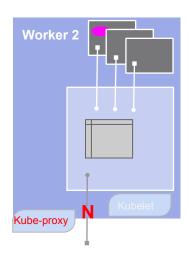


### The Kubernetes Approach: Services



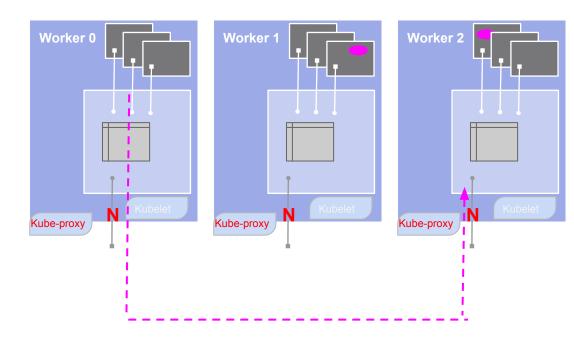






### The Kubernetes Approach: Services Cluster IP





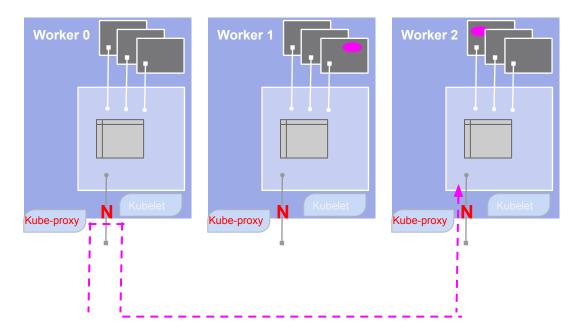
### The Kubernetes Approach: Services

#### **Node Port**

```
apiVersion: v1
kind: Service
metadata:
  name: nginx-service
labels:
    app: nginx
spec:
  type: NodePort
ports:
    - port: 80
    nodePort: 30080
selector:
```

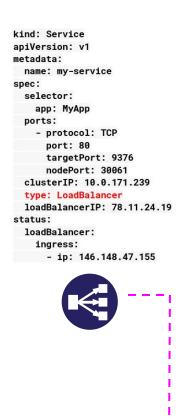
app: nginx

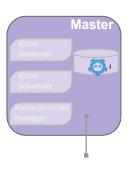


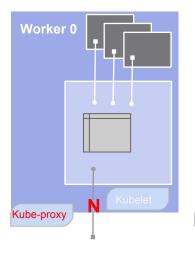


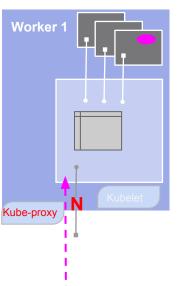
### The Kubernetes Approach: Services

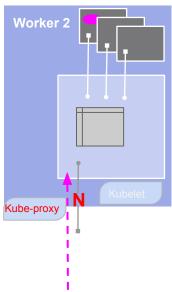
### **External Load Balancers**





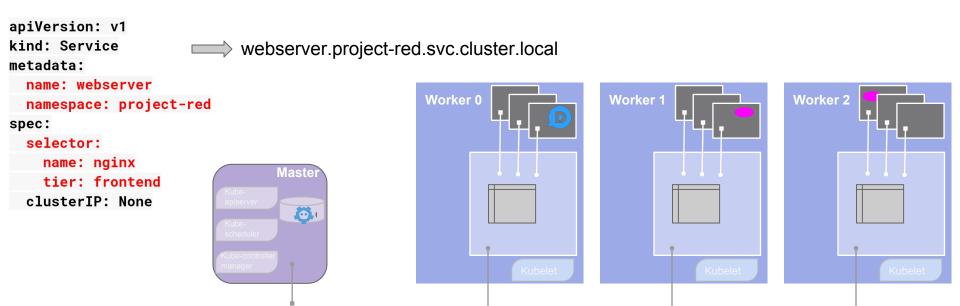






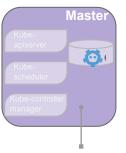
### The Kubernetes Approach: KubeDNS

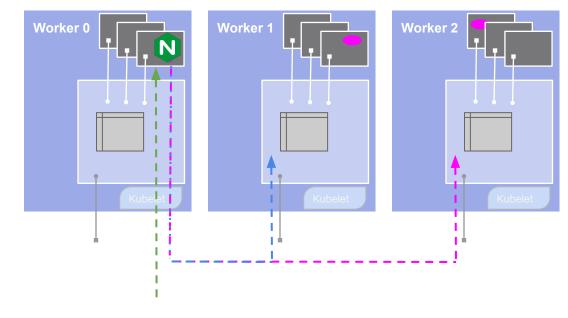
### Naming and Service Discovery



### The Kubernetes Approach: Ingress Resources & Controllers

apiVersion: extensions/v1beta1 kind: Ingress metadata: name: test spec: rules: - host: foo.bar.com http: paths: - backend: serviceName: s1 servicePort: 80 - host: bar.foo.com http: paths: - backend: serviceName: s2 servicePort: 80



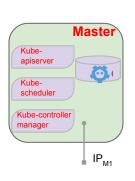


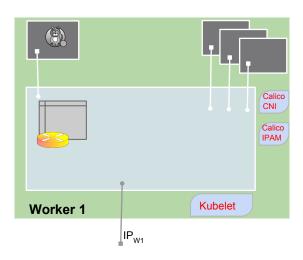
#### Demo: Self-hosted Kube+Calico Cluster with kubeadm

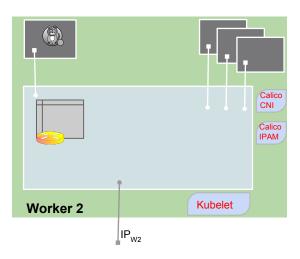
Step 1: *kubeadm init* (on master)

Step 2: kubectl apply -f calico.yaml

Step 3: *kubeadm join* (on each worker)









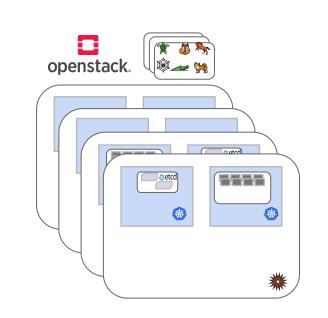
### Common Network Operations Across K8s & OpenStack

Tomorrow (Wed): Hynes Level 3 - Ballroom B - 4:30pm



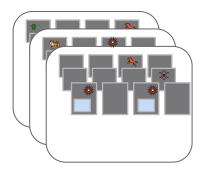


Kubernetes alongside OpenStack



Kubernetes in OpenStack





OpenStack in Kubernetes

### **Takeaways**



Networking at VM-scale != Networking at Container Scale

### **Takeaways**



Kubernetes Network Abstractions are different from Neutron

### **Takeaways**



Keep network simple/scalable, Network Policy for isolation